

[CS6240] Large-Scale Parallel Data Processing

Homework03

Hwijong Im

Problem01

(1). RDD-G Pseudo Code

```
val conf = new SparkConf().setAppName("RDD-G").setMaster("local")
val sc = new SparkContext(conf)
```

```
val lines = sc.textFile("edges.csv")
val pairs = lines.map(line => {
    val parts = line.split(",")
    (parts(1).toInt, 1)
}).filter(_._1 % 100 == 0)
```

```
val grouped = pairs.groupByKey()
val counts = grouped.mapValues(_.sum)
```

```
counts.saveAsTextFile("output/RDD-G")
sc.stop()
```

(2). RDD-R Pseudo Code

```
val conf = new SparkConf().setAppName("RDD-R").setMaster("local")
val sc = new SparkContext(conf)
```

```
val lines = sc.textFile("edges.csv")
val pairs = lines.map(line => {
    val parts = line.split(",")
    (parts(1).toInt, 1)
}).filter(_._1 % 100 == 0)
```

```
val counts = pairs.reduceByKey(_ + _)
```

```
counts.saveAsTextFile("output/RDD-R")
sc.stop()
```

(3). RDD-F Pseudo Code

```
val conf = new SparkConf().setAppName("RDD-F").setMaster("local")
val sc = new SparkContext(conf)
```

```
val lines = sc.textFile("edges.csv")
```

```

val pairs = lines.map(line => {
  val parts = line.split(",")
  (parts(1).toInt, 1)
}).filter(_._1 % 100 == 0)

val counts = pairs.foldByKey(0)(_ + _)

counts.saveAsTextFile("output/RDD-F")
sc.stop()

```

(4). RDD-A Pseudo Code

```

val conf = new SparkConf().setAppName("RDD-A").setMaster("local")
val sc = new SparkContext(conf)

val lines = sc.textFile("edges.csv")
val pairs = lines.map(line => {
  val parts = line.split(",")
  (parts(1).toInt, 1)
}).filter(_._1 % 100 == 0)

val zeroValue = 0
val counts = pairs.aggregateByKey(zeroValue)(
  (acc, value) => acc + value,
  (acc1, acc2) => acc1 + acc2
)

counts.saveAsTextFile("output/RDD-A")
sc.stop()

```

(5). DSET Pseudo Code

```

val spark = SparkSession.builder()
  .appName("DSET")
  .master("local")
  .getOrCreate()

import spark.implicits._

val df = spark.read.csv("edges.csv")
  .toDF("user", "follower")
  .select($"follower".cast("int").as("follower"))
  .filter($"follower" % 100 === 0)
  .groupBy($"follower")
  .agg(count($"follower").as("num_followers"))

df.write.format("csv").save("output/DSET")
spark.stop()

```

Problem02

Code Link of Aggregations

RDD-G:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/spark_g/RDD_G.java

RDD-R:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/spark_r/RDD_R.java

RDD-F:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/spark_f/RDD_F.java

RDD-A:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/spark_a/RDD_A.java

DSET:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/spark_dset/RDD_DS ET.java

Problem03

(1). RDD-G

The shuffling has been proceeded before using GroupByKey.

```
24/02/16 21:16:29 INFO SparkContext: Created broadcast 0 from textFile at RDD_G.java:54
24/02/16 21:16:29 INFO FileInputFormat: Total input files to process : 1
(40) MapPartitionsRDD[6] at mapValues at RDD_G.java:68 []
| MapPartitionsRDD[5] at groupByKey at RDD_G.java:67 []
| ShuffledRDD[4] at groupByKey at RDD_G.java:67 []
+- (40) MapPartitionsRDD[3] at filter at RDD_G.java:64 []
| MapPartitionsRDD[2] at mapToPair at RDD_G.java:58 []
| input/edges.csv MapPartitionsRDD[1] at textFile at RDD_G.java:54 []
| input/edges.csv HadoopRDD[0] at textFile at RDD_G.java:54 []
```

(2). RDD-R

ReduceByKey operation includes a shuffling stage where local aggregation is performed before data is shuffled across the network. the “ShuffledRDD[4]” indicates that shuffling occurs as part of the reduceByKey operation.

```
Execution Plan for followerCounts:
(40) ShuffledRDD[4] at reduceByKey at RDD_R.java:56 []
+- (40) MapPartitionsRDD[3] at filter at RDD_R.java:53 []
| MapPartitionsRDD[2] at mapToPair at RDD_R.java:47 []
| input/edges.csv MapPartitionsRDD[1] at textFile at RDD_R.java:43 []
| input/edges.csv HadoopRDD[0] at textFile at RDD_R.java:43 []
24/02/16 21:50:22 INFO SparkContext: Starting job: foreach at RDD_R.java:66
24/02/16 21:50:22 INFO DAGScheduler: Registering RDD 3 (filter at RDD_R.java:53) as input to shuffle 0
24/02/16 21:50:22 INFO DAGScheduler: Got job 0 (foreach at RDD_R.java:66) with 40 output partitions
24/02/16 21:50:22 INFO DAGScheduler: Final stage: ResultStage 1 (foreach at RDD_R.java:66)
24/02/16 21:50:22 INFO DAGScheduler: Parents of final stage: List(ShuffleMapStage 0)
24/02/16 21:50:22 INFO DAGScheduler: Missing parents: List(ShuffleMapStage 0)
```

(3). RDD-F

FoldbyKey operation also includes a shuffling stage to aggregate effectively.

```
Execution Plan for FoldByKey:
(40) ShuffledRDD[4] at foldByKey at RDD_F.java:55 []
+- (40) MapPartitionsRDD[3] at filter at RDD_F.java:52 []
| MapPartitionsRDD[2] at mapToPair at RDD_F.java:46 []
| input/edges.csv MapPartitionsRDD[1] at textFile at RDD_F.java:42 []
| input/edges.csv HadoopRDD[0] at textFile at RDD_F.java:42 []
```

(4). RDD-A

AggregatebyKey operation starts at the nearly same with shuffling stage.

```

Execution Plan for Aggregate:
(40) ShuffledRDD[4] at aggregateByKey at RDD_A.java:57 []
+- (40) MapPartitionsRDD[3] at filter at RDD_A.java:53 []
    | MapPartitionsRDD[2] at mapToPair at RDD_A.java:47 []
    | input/edges.csv MapPartitionsRDD[1] at textFile at RDD_A.java:43 []
    | input/edges.csv HadoopRDD[0] at textFile at RDD_A.java:43 []

```

(5). DSET

As you can see the image, the aggregation operation, groupby, has been performed after shuffling stage. This indicates that Spark collects and redistributes the data based on the grouping key("userid") through shuffling, and applies the aggregation function "count" to compute the final results.

```

== Physical Plan ==
AdaptiveSparkPlan isFinalPlan=false
+- HashAggregate(keys=[userId#21], functions=[count(followerId#22)])
   +- Exchange hashpartitioning(userId#21, 200), ENSURE_REQUIREMENTS, [plan_id=34]
      +- HashAggregate(keys=[userId#21], functions=[partial_count(followerId#22)])
         +- Project [_c0#17 AS userId#21, _c1#18 AS followerId#22]
            +- Filter (isNotNull(_c0#17) AND ((_c0#17 % 100) = 0))
               +- FileScan csv [_c0#17, _c1#18] Batched: false, DataFilters: [isNotNull(_c0#17), ((_c0#17 % 100) = 0)],
Format: CSV, Location: InMemoryFileIndex(1 paths)[file:/Users/minsungim/Desktop/Desktop/hw-3-bellwood22-master/input/
edg..., PartitionFilters: [], PushedFilters: [IsNotNull(_c0)], ReadSchema: struct<_c0:int,_c1:int>

== Parsed Logical Plan ==
'Aggregate [userId#21], [userId#21, count('followerId) AS followercount#30]
+- Filter ((userId#21 % 100) = 0)
   +- Project [_c0#17 AS userId#21, _c1#18 AS followerId#22]
      +- Relation [_c0#17, _c1#18] csv

== Analyzed Logical Plan ==
userId: int, followercount: bigint
Aggregate [userId#21], [userId#21, count(followerId#22) AS followercount#30L]
+- Filter ((userId#21 % 100) = 0)
   +- Project [_c0#17 AS userId#21, _c1#18 AS followerId#22]
      +- Relation [_c0#17, _c1#18] csv

== Optimized Logical Plan ==
Aggregate [userId#21], [userId#21, count(followerId#22) AS followercount#30L]
+- Project [_c0#17 AS userId#21, _c1#18 AS followerId#22]
   +- Filter (isNotNull(_c0#17) AND ((_c0#17 % 100) = 0))
      +- Relation [_c0#17, _c1#18] csv

== Physical Plan ==
AdaptiveSparkPlan isFinalPlan=false
+- HashAggregate(keys=[userId#21], functions=[count(followerId#22)], output=[userId#21, followercount#30L])
   +- Exchange hashpartitioning(userId#21, 200), ENSURE_REQUIREMENTS, [plan_id=34]
      +- HashAggregate(keys=[userId#21], functions=[partial_count(followerId#22)], output=[userId#21, count#34L])
         +- Project [_c0#17 AS userId#21, _c1#18 AS followerId#22]
            +- Filter (isNotNull(_c0#17) AND ((_c0#17 % 100) = 0))
               +- FileScan csv [_c0#17, _c1#18] Batched: false, DataFilters: [isNotNull(_c0#17), ((_c0#17 % 100) = 0)],
Format: CSV, Location: InMemoryFileIndex(1 paths)[file:/Users/minsungim/Desktop/Desktop/hw-3-bellwood22-master/input/
edg..., PartitionFilters: [], PushedFilters: [IsNotNull(_c0)], ReadSchema: struct<_c0:int,_c1:int>

```

Problem04

(1). RS-R

```

object RS_R_Triangles {
  def main(args: Array[String]): Unit = {
    val conf = new SparkConf().setAppName("RS_R_Triangles").setMaster("local")
    val sc = new SparkContext(conf)

    // Load data
    val rawData: RDD[String] = sc.textFile("input/edges.csv")

    // Create (a, b) and (b, a) pair RDDs
    val edges: RDD[(String, String)] = rawData.flatMap(line => {

```

```

    val parts = line.split(",")
    Seq((parts(0), parts(1)), (parts(1), parts(0)))
  })

  // Create pairs for potential triangles (a, b) => (b, a)
  val possibleTriangles: RDD[((String, String), String)] = edges.map {
    case (a, b) => ((a, b), b)
  }

  // Join on (a, b) to find all potential triangles
  val joined: RDD[((String, String), ((String, String), String))] =
possibleTriangles.join(possibleTriangles)

  // Filter to find actual triangles
  // This step filters out false positives by ensuring that the third edge exists to close the
triangle
  val triangles: RDD[(String, String, String)] = joined.flatMap {
    case ((a, b), ((_, _), c)) if a != c && b != c =>
      Seq((a, b, c))
    case _ => Seq.empty[(String, String, String)]
  }.distinct()

  // Count the triangles
  val triangleCount: Long = triangles.count()

  println(s"Total Triangles: $triangleCount")
  triangles.saveAsTextFile("output/RS_R_Triangles")

  sc.stop()
}
}

```

(2). RS-D

```

object RS_D {
  def main(args: Array[String]): Unit = {
    // Spark setting
    val spark = SparkSession.builder.appName("RS_D").master("local").getOrCreate()

    // Dataset import
    val edges = spark.read.option("inferSchema", "true").option("header", "false")
      .csv("input/edges.csv").toDF("src", "dst")
    val reversedEdges = edges.select(col("dst").alias("src2"), col("src").alias("dst2"))

    // join_01: Find pairs of edges that could potentially form triangles
    val joined = edges.as("edges1").join(reversedEdges.as("edges2"),
      col("edges1.dst") === col("edges2.src2"))

    // join_02: Complete triangles by finding the third edge

```

```

val triangles = joined.as("joined")
    .join(edges.as("edges3"), col("joined.dst2") === col("edges3.src"))
    .where(col("joined.src") === col("edges3.dst"))

// Output
triangles.show()
val trianglesCount = triangles.count()
println(s"Total count: $trianglesCount")

// Save output
triangles.write.format("csv").option("header",
"true").mode("overwrite").save("output/RS_D/triangles")

spark.stop()
}

```

(3). REP-R

```

object REP_R {
  def main(args: Array[String]): Unit = {
    // Spark setting
    val conf = new SparkConf().setAppName("REP_R").setMaster("local")
    val sc = new SparkContext(conf)

    // data load
    val rawData: RDD[String] = sc.textFile("input/edges.csv")

    // (a, b), (b, a) pair
    val df01: RDD[(String, String)] = rawData.map(line => {
      val parts = line.split(",")
      (parts(0), parts(1))
    })
    val df02: RDD[(String, String)] = rawData.map(line => {
      val parts = line.split(",")
      (parts(1), parts(0))
    })

    // df02 mapping and broadcasting
    val broad_df02: Broadcast[Map[String, String]] = sc.broadcast(df02.collectAsMap())

    // perform join
    val triangles: RDD[String] = df01.flatMap {
      case (a, b) =>
        val cOption = broad_df02.value.get(b)
        if (cOption.isDefined && broad_df02.value.contains(cOption.get) &&
broad_df02.value(cOption.get) == a) {
          Some(s"$a,$b,$cOption.get")
        } else {
          None
        }
    }
  }
}

```

```

    }
}

// triangle counts
val triangleCount: Long = triangles.count()
println(s"Total Triangles: $triangleCount")

// result
triangles.saveAsTextFile("output/REP_R")

sc.stop()
}
}

```

(4). REP-D

```

object REP_D {
  def main(args: Array[String]): Unit = {
    // Spark setting
    val spark = SparkSession.builder()
      .appName("REP_D")
      .master("local")
      .getOrCreate()

    import spark.implicits._

    // data load
    val edges: Dataset[Row] = spark.read
      .option("inferSchema", "true")
      .csv("input/edges.csv")
      .toDF("src", "dst")

    // create opposite dataset
    val reversedEdges = edges.select($"dst".alias("b"), $"src".alias("a"))

    // broadcasting
    val broadcastedReversedEdges = broadcast(reversedEdges)

    // first join
    val join01 = edges.join(broadcastedReversedEdges, $"dst" === $"b")
      .select($"src".alias("a"), $"dst".alias("b"), $"a".alias("c"))

    // second join
    val triangles = join01.join(edges, $"c" === edges("dst") && $"a" === edges("src"))
      .select($"a", $"b", $"c")

    // output
    triangles.show(false)
    val triangleCount = triangles.count()
  }
}

```

```
println(s"Total Triangles: $triangleCount")

// save
triangles.write
  .format("csv")
  .option("header", "true")
  .mode("overwrite")
  .save("output/REP_D/triangles")

spark.stop()
}
}
```

Problem05

Github Link of Join Code

RS-R:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/join_rs_r/RS_R.java

RS-D:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/join_rs_d/RS_D.java

REP-R:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/join_rep_r/REP_R.java

REP-D:

https://github.com/CS6240/hw-3-bellwood22/blob/master/src/main/java/join_rep_d/REP_D.java

Problem06

(1). RS-R

```
24/02/17 22:44:09 WARN SparkConf: The configuration key 'spark.yarn.executor.memoryOverhead' has been deprecated as of Spark 2.3 and may be removed in the future. Please use the new key 'spark.executor.memoryOverhead' instead.
Warning: Skip remote jar S3://mr-median-hw1j0ng/spark-demo.jar.
24/02/17 22:44:11 INFO RMPProxy: Connecting to ResourceManager at ip-172-31-88-101.ec2.internal/172.31.80.101:8032
24/02/17 22:44:11 INFO Client: Requesting a new application from cluster with 4 NodeManagers
24/02/17 22:44:11 INFO Client: Verifying our application has not requested more than the maximum memory capability of the cluster (11520 MB per container)
Exception in thread "main" java.lang.IllegalArgumentException: Required executor memory (16240+2048 MB) is above the max threshold (11520 MB) of this cluster! Please check the values of 'yarn.scheduler.maximum-allocation-mb' and/or 'yarn.nodemanager.resource.memory-mb'.
    at org.apache.spark.deploy.yarn.Client.verifyClusterResources(Client.scala:318)
    at org.apache.spark.deploy.yarn.Client.submitApplication(Client.scala:166)
    at org.apache.spark.deploy.yarn.Client.run(Client.scala:1152)
    at org.apache.spark.deploy.yarn.YarnClusterApplication.start(Client.scala:1520)
    at org.apache.spark.deploy.SparkSubmit$.org$apache$spark$deploy$SparkSubmit$$runMain(SparkSubmit.scala:894)
    at org.apache.spark.deploy.SparkSubmit$.doRunMain$1(SparkSubmit.scala:198)
    at org.apache.spark.deploy.SparkSubmit$.submit(SparkSubmit.scala:228)
    at org.apache.spark.deploy.SparkSubmit$.main(SparkSubmit.scala:137)
    at org.apache.spark.deploy.SparkSubmit.main(SparkSubmit.scala)
24/02/17 22:44:11 INFO ShutdownHookManager: Shutdown hook called
24/02/17 22:44:11 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-efa22369-8b68-44c6-be3a-813b8adce3df
Command exiting with ret '1'
```

Tried to solve the memory issue, however, even when I tried to extend memory, the same issue happened repeatedly. I lowered the max value of executor memory to 8GB with a memory overhead of 3GB.

machine type: 1 master node, 4 worker node

max value: 8 driver memory, 8 executor memory, 3 of each memory overhead

running time: NA

triangle count: NA

(2). RS_D

machine type: 1 master node, 4 worker node
max value: 8 driver memory, 8 executor memory, 3 of each memory overhead
running time: NA
triangle count: NA

```
client token: N/A
diagnostics: Application application_1708296077248_0001 failed 2 times due to AM Container for appatempt_1708296077248_0000002 exited with  exitCode: 137
Failing this attempt. Diagnostics: Container killed on request. Exit code is 137
Container exited with a non-zero exit code 137
Killed by external signal
To check the application tracking page: http://ip-172-31-94-43.ec2.internal:8888/cluster/app/application\_1708296077248\_0001 Then click on Links to logs of each attempt.
. Failing the application.
ApplicationMaster host: N/A
ApplicationMaster RPC port: -1
queue: default
start time: 1708296121677
final status: FAILED
```

machine type: 1 master node, 4 worker node
max value: 8 driver memory, 8 executor memory, 3 of each memory overhead
running time: NA
triangle count: NA

machine type: 1 master node, 4 worker node
max value: 8 driver memory, 8 executor memory, 3 of each memory overhead
running time: 143343 milliseconds
triangle count: 2336

https://mr-median-hwijong.s3.amazonaws.com/log/j-33DCSHZPY8INN/steps/s-0559237WG
G845J0CYCR/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIA
WVEGCM7SYTVDMDWD%2F20240218%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-D
ate=20240218T204016Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjE
B0aCXVzLWVhc3QtMSJHMEUCIBix1QG%2F8PahhlEoz%2FUm0orwLxhy9Dj1L2V46zKOls
sMAiEAmdcD5yn7z9yID%2F44rRLglEjAH7zp0hQ44bWXzfbHdUEqiqMI9v%2F%2F%2F%2
F%2F%2F%2F%2F%2F%2F%2FARAAGw0NTc2OTUxOTMwNjEiDldhK57MRSCBfPC4JyreAv
bl%2FHzHN54Z39gHVt8A5RBZKI0pRWBm2fTsRIWbB0VM67VFuSxVcp%2BERGeaBtdOFv
KYx9leGcsfPgwrwStnl8XMP9spLcizJdK6zFcmf3uogT3fmfUCmG3gBWcMWxt%2F4K2n5G
wd5tu3R1PNOTUZEHFknhkuQEmHczvJ65hRfd88wrtcrhS2VXjm1j14Q%2BvjWrlI4ZVBd1JM

X0BnPCXJVAXN%2FuUumtPMguKU2FyT8%2B7K0uI0knAWi%2BcAxGR6DoNXN25VNWRsgjw%2FDq1MKcAbUZU6KI3iXW3e9sG39X9nWqWsiyhjrcliJyq%2BSheMswnlZDLdM8LQ4FHpK6xHZI9SxDcMf1nm8I96UtUPAXeA37XgmwWsrLxIUcZvPhRrJOpZ%2Bk0fCNlvp3GgL7Ded%2FY4773%2FjJC1ni3jvWvdBoQEvn92uubwLpZMA8HLnWq3EwlSyqoctsMxNBD2Q%2Bk%2BBX8MLm8ya4GOocCGfAlkZJ33%2BbkvOhYkTkkgFhrZYI09PCtvxyLT1u%2BaeXNVsoDuwNi0YLoCQ8e1aTMocUdKCLHkZ3pRzcqsBbKtfgIII%2BQBLT01H6m09cSmWtDFUJZqYHzLggCZYRfpYw54ud9L4iogm4w%2FBB5x890ILhl8zIIM%2BxcLpanB5tFNQbMSOP91J6juCl6kXlztDbkR27auD1ClqCRnOJtJlZHO540TAtx%2BziuLGdE2EYsBk%2Bd88QRZaReU7OPASHd8KQIPEYqInVg7r3trxeNxJN%2F7Su%2F2nu1kNt26Yb2TdZC7IwV9LQzc7IzXUjZGq86d5X36fxWIVKp2E2bUADk6sJURHjS2qEacY%3D&X-Amz-Signature=ee9068443750d79b6001aee2261ff3252b3fb67c0b8b42c55bfc7cf9b0b0b9&X-Amz-SignedHeaders=host&response-content-disposition=inline

(2). RS_D Link:

[https://mr-median-hwijong.s3.amazonaws.com/log/j-3MH5ZMSM4EAVF/steps/s-07210232E5WS81GL3XBR/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIAWVEGCM7SWBYO2GX6%2F20240218%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20240218T211009Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEB0aCXVzLWVhc3QtMSJGMEQCID53DY%2BbiBOYUH8jBOVw2qSma8f3gdg4bywm9zc8MK2TAiBdEvrsYHY6DsufHvOMOo1wkrnosij7IA%2FYNSDXcyhJvyqKAwj2%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F8BEAAaDDQ1NzY5NTE5MzA2MSIM9sDVnA2G6Y7ng0aOKt4CXW0gSkq4Px59dJ%2BlaPJ4vL1USilRptmXhj88JyP7KgmBQZfx%2BY8CZyzqZGsYNxoNujRGDOVthW3VWmi%2BIPm7MlImUCqp0nyG0sAo8LB7MLRwb8KSPzkxgyksl3kjYu79v5%2FwoF0WT7IxnIryZT5IQF4VKG2O3snk7R75nybEB2TF5Z5w5xfWoQ4PLm0sullSxNW81pKFJLj34rfU6%2FcZyI%2F208Y9KENZJAPHuIUpZVXzaM6ogCuREhrStKJxNUP9MQqEwgZ9heZWHZjfy4b01tX1pVAOEscqhivKyes%2Fn74qJ4Zt4FzlyTcRvZoqjT%2B2ny6Qeul5r0MDWx92FXTiGp%2F7bjHlwF0woWIYApe%2Fw8v0eJ%2BosD%2FiiMm8AdKoKIT%2FshAu%2Bg37QUora8NecpLa2bSj9Sa%2BMOVfYpfr1HZ9S1NfKGv9Mf544t5jaY5gFYGLAQXATemql0%2FmDBEiKEI0wubzJrgY6iAL8b5S3ZuEQInnMmjFEokRfRhRxKes2R%2BHmRML3shfjdhRvhOPmvdPVNsZWc8QAmdABxOizk3DqxQsX0Xb0QW5xQCmAIVnMt0WD3V%2FY2zRfPowYRutGB%2B7PwaomHZ997OC8XW%2FN4LdD6Vr3Lg9gjvd52t5AdqC4BeEeQgGNkvVXFdi51m2m8PxpGziSSPoREA4dM17vWcl2h0A7%2Bexh06VoBz%2BSqG4liW4ogunuTzAShe%2B4hwjzOpVu%2B395u%2F%2B66eISTE%2F9uSo3junyQEkTmlP9vwDwUaeMsl6oGsmZzneEzIGFRxaaY4k4njFsMZvJHh9nF3s1UGOF6BCEZRsx%2FQNSofcllrtBu1o%3D&X-Amz-Signature=1ff102dd3cf2e43b53ca9f862672ef104c8a4eeb367014e72c4d1381990edf3e&X-Amz-SignedHeaders=host&response-content-disposition=inline](https://mr-median-hwijong.s3.amazonaws.com/log/j-3MH5ZMSM4EAVF/steps/s-07210232E5WS81GL3XBR/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIAWVEGCM7SWBYO2GX6%2F20240218%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20240218T211009Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEB0aCXVzLWVhc3QtMSJGMEQCID53DY%2BbiBOYUH8jBOVw2qSma8f3gdg4bywm9zc8MK2TAiBdEvrsYHY6DsufHvOMOo1wkrnosij7IA%2FYNSDXcyhJvyqKAwj2%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F8BEAAaDDQ1NzY5NTE5MzA2MSIM9sDVnA2G6Y7ng0aOKt4CXW0gSkq4Px59dJ%2BlaPJ4vL1USilRptmXhj88JyP7KgmBQZfx%2BY8CZyzqZGsYNxoNujRGDOVthW3VWmi%2BIPm7MlImUCqp0nyG0sAo8LB7MLRwb8KSPzkxgyksl3kjYu79v5%2FwoF0WT7IxnIryZT5IQF4VKG2O3snk7R75nybEB2TF5Z5w5xfWoQ4PLm0sullSxNW81pKFJLj34rfU6%2FcZyI%2F208Y9KENZJAPHuIUpZVXzaM6ogCuREhrStKJxNUP9MQqEwgZ9heZWHZjfy4b01tX1pVAOEscqhivKyes%2Fn74qJ4Zt4FzlyTcRvZoqjT%2B2ny6Qeul5r0MDWx92FXTiGp%2F7bjHlwF0woWIYApe%2Fw8v0eJ%2BosD%2FiiMm8AdKoKIT%2FshAu%2Bg37QUora8NecpLa2bSj9Sa%2BMOVfYpfr1HZ9S1NfKGv9Mf544t5jaY5gFYGLAQXATemql0%2FmDBEiKEI0wubzJrgY6iAL8b5S3ZuEQInnMmjFEokRfRhRxKes2R%2BHmRML3shfjdhRvhOPmvdPVNsZWc8QAmdABxOizk3DqxQsX0Xb0QW5xQCmAIVnMt0WD3V%2FY2zRfPowYRutGB%2B7PwaomHZ997OC8XW%2FN4LdD6Vr3Lg9gjvd52t5AdqC4BeEeQgGNkvVXFdi51m2m8PxpGziSSPoREA4dM17vWcl2h0A7%2Bexh06VoBz%2BSqG4liW4ogunuTzAShe%2B4hwjzOpVu%2B395u%2F%2B66eISTE%2F9uSo3junyQEkTmlP9vwDwUaeMsl6oGsmZzneEzIGFRxaaY4k4njFsMZvJHh9nF3s1UGOF6BCEZRsx%2FQNSofcllrtBu1o%3D&X-Amz-Signature=1ff102dd3cf2e43b53ca9f862672ef104c8a4eeb367014e72c4d1381990edf3e&X-Amz-SignedHeaders=host&response-content-disposition=inline)

(3). REP_R Link:

https://mr-median-hwijong.s3.amazonaws.com/log/j-3A1HL75C67P4R/steps/s-02438524O6A90GGESBV/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIAWVEGCM7SSOADV3HI%2F20240218%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20240218T225029Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEB8aCXVzLWVhc3QtMSJGMEQCIBWiralhPBCGkuEottcagTrVjwZSkZICAeq%2BXxoQO2eXAiBZqefAPtAiJDq%2Ft%2FFh%2BKGWWhHruW94KN%2FgnTQHIDVM90yqKAwj4%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F8BEAAaDDQ1NzY5NTE5MzA2MSIMadb26wbSh3PlycsVKt4CPPfCWu6ujBOobrDvVk%2B5iy2EccdvbxJAQR49s286Dcd04Vs%2Fz0s6Dpl0gfUB9J3be%2B1W5w%2Bkpdzj6AIQBuiS%2B6qvI3ZYUYkBJ2rtFXU1SvxcXb%2BcJ268Fg1tUuZTDphzHUuVa1z3OA5Inum35MsQx8yNaq2hzLQgWc3FgDYb0cYuOjAqO5cbJhL4vTU2d7jU9Ve

(4). REP D Link:

Problem08

Nevertheless ran 2 hours to perform the program, it did not completed but only keep running.

max value: 8 driver memory, 8 executor memory, 3 of each memory overhead

triangle count: na

Nevertheless ran 2 hours to perform the program, it did not completed but only keep running. Because of the join inefficiency, the program code is hard to get the result of the number of triangles.

triangle count: na

triangle count: na

triangle count: 135588 milliseconds

https://mr-median-hwijong.s3.amazonaws.com/log/j-2CHI9HQXOYUM2/containers/application_1708297787674_0001/container_1708297787674_0001_01_000001/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIAWVEGCM7STPLDEEOE%2F20240218%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20240218T233318Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjECAaCXVzLWVhc3QtMSJHMEUCIGqb8d2sZ9m%2BNoGXdQ7CETyKE7GOkI%2BS5aUKKRXgYmtFAiEAlXuOHWXkF4%2BwrNvXlaAaqbs7A8S2LWhzgs6VLvhqqyEqigMI%2Bf%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F%2FARAAGw0NTc2OTUxOTMwNjEiDDLvIkJH6G%2B%2BorPqlireAprmsHNjntYW4R5zU9%2BLmI%2Fd4YXO68yUYw3mynoQp9%2B2AoHDRT1w558bH7Ww8B12aZfGljgKAX0bEjLjuJiL00I5cBtf5OzgTilzilmEkax6tqcl657Cn9JIZRXUFxwXIsWs3y7%2B2bQ92%2B3I95WsNQQnZ%2FlahAaVh9W2%2FQxhFwfYBQ23hEmrsmzy7Wi1UyrRqexmUZ3gFgddledQWce99ajlvSHqXLC%2BoQil8K9c8iJ%2By9Q4FVGskqQRu0WiKbQb1%2FugncEBaUhS5FyqhFbXu4%2FUeuU9rrt3wspk%2Femjn7UocR9nYdBX43T9k9PSz7PvYm1%2BvjLDgV3rbBIO9BVhE3y2PbqRiTzMEC9zk5NZz4u8LWbWal9UypgD9TAPdBe18jtdESBk%2FsPlc1jdHH5gcCScO0walb1GwbvtwE6RTUHLekvaA4p0%2BLbhE%2BDPK1yf642IWfdX%2FNVnCS%2FNbKQTMLm8ya4GOocC8xpTu4kx6S6LKDDoan%2FvyUH%2BNQ%2FsZh8w8DibsFDPRnV%2BdsKCxs%2Fs6FvBKanzCuB1hzJYoyip7ucvRppJUS0i7M1neWLS%2FFF8m7IRZFCNP5o4hNP5LF2mF83k%2Fb0pRubGMxYuB6Fa997kuJimUhEclPerDn3oOw3lx3lrloz7HHjM5fcJsru

RS-D Link:

REP-R:

https://mr-median-hwijong.s3.amazonaws.com/log/j-D9UWLDORVFAY/steps/s-00285333916
G9OH5CJZZ/stderr.gz?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIA
WVEGCM7SWNRAVQXJ%2F20240219%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Da
te=20240219T005318Z&X-Amz-Expires=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEC
EaCXVzLWVhc3QtMSJIMEYCIQCCtM8fE1nvCb%2BpPhUd5%2Frb8ymS1bnUUXJwZA9e
KglfQlhAPURTMxfFYsVEZ7YNQrXoH%2BCCH6p2%2BJV8h4uoVdeFimKooDCPr%2F%2
F%2F%2F%2F%2F%2F%2F%2F%2FwEQABoMNDU3Njk1MTkzMDYxl9x2vpo7bHBbjf%2B
J0GAq3gJMw5YsMUGqx3gXClsiOlqTjX10Uo2kn%2Fgf52HtFtJt8t9sWRaA2GxRBZIX4gLC
7EZ30RdxD6CSWvfoX%2B2DK9%2Bz62eztOJVn3yeQkpC4SmW11wINVYOnrIWJnVgJ%
2FymRBKjVn9X2euW1uAZcmGHPtvou0FiwkwXyYnYNFdOnq%2Fg8OEKxJN0LSEfampzmJ
E2D6kX9Q7nwtj5c5KVxKUK3oGC7iulB6KaLOvb%2FCWK%2B0QWaz2tjcKYAZV69pb1RrC
D6LDm6%2F4%2FleN%2B53UFHl5qTjPY5htfkeUzKVIZW26CrFCPgxDxFcX8OYPlv896rEj
YgH%2F%2BhbRud3RGYWj7HVyax%2FJkK3M2syNEKCJIYbS7uapPNfLwT%2B%2B6Es
W6TH%2FBExLIQxS2XEWOnrKKfDxE0J9lbzTDNiXywflDNGFKvQO7RwGxT3sMIlx7xlbP
C1CGdbLzSAGILmclvoCLBVtmQDxCjCOssquBjqGAR8pG8hfcuUtpQsqA06Pvt1JrUSQvmw
kBO8%2FQAFrUeOOSgDiygeXZAFstNPEG0iL4szAwk6afoXm3TdoqLIHTBTHW51yu93QBB
8i3kKm3BijMZE8BTpqwcnFLrhrQQumhKrLSq%2B%2BfrMxQFZs3s5oiANLCTtbrBa7rmtliVC

PRml8HxO%2Bo6RWc13jRPOv1fwjcv8zEt%2FiSRp32AoYv94ZHenHlcUJhqZxHYvyl98bU
%2BZekAyTn76y%2FDwjJSdHGI0CRWJPN3mk31wPHKizTC%2Bk9OHdtJYPm1j7HX7dHtr
5cw2SU0lzZEMFrD8ytxKCrtnK7B%2F7AkkbTs9IlbqOMHZgg4RftMKI4Zo%3D&X-Amz-Sign
ature=9e44baa835f14e72f931e4c3f532ebaa8843e12beaec241f5c81d324855ab762&X-Amz
-SignedHeaders=host&response-content-disposition=inline

REP-D:

https://mr-median-hwijong.s3.amazonaws.com/log/j-F1JXU1D1JSKI/containers/application_
1708304365290_0001/container_1708304365290_0001_01_000001/stderr.gz?X-Amz-Algor
ithm=AWS4-HMAC-SHA256&X-Amz-Credential=ASIAWVEGCM7S66VSZVPG%2F2024021
9%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20240219T010814Z&X-Amz-Expire
s=900&X-Amz-Security-Token=IQoJb3JpZ2luX2VjEClaCXVzLWVhc3QtMSJGMEQCIaA8lw
LrK09tCgBK61bo4DklXymCSsH3jcKH0PH4JwfwAiA33uZszse374URDdDwLI7x0us8RDvn4
QJ0x3VS8GF61yqKAwj6%2F%2F%2F%2F%2F%2F%2F%2F%2F%2F8BEAAaDDQ1NzY5
NTE5MzA2MSIMlnZY8Xgvm4zBqp5gKt4ChpxF76DsAUsvIFRTjwltikeaUvwm7Rh9ML36Qv
PrSI%2BLTqotYgXEO%2BqQA8vFgOV1gh4CeYWCQMHsBi8ACytgX%2BYC7mBopuygb
FJpY7kHKtfOJhCokxMDL1ejoe2BiGBXA%2BFngy3Lx%2FMmebMMkgFrY%2F4UH8NPrm
7EkIRM9ek%2FxDyzIkKj%2FE8xTEQHAOWVipXQ6BrsVa0R51K9dMpRd8k4v%2BerlsNek
SxhagVqnYaZrmy059uRdtjSvblur%2B4njbuxLzbEjqJMbSpdn2Q5dmUu8D00XyKKOaEd5G
wGeQgkl5eFeedsZhJD%2FRM6j3EXuQiURakp%2FD0ZHTFGePCJ6Y7lQtA7oz4VFf%2F2x
yQyXWyRi4uD0nTmd25DnYICUclsUevS%2BdlugS0rkTCLkbN%2BnURKgcGA7CqqFW97l
LBPilrQCpDE%2FZpe5Vy1VtAO8j7r%2FKIMndr3Ge%2ByaTluYYLswjrLKrgY6iALoOzCEZU
2b1PeO5rSGEUlvpXLldIKyUa447oweozdFUcvl6qsnOTMeIYma4lWo8mNXGje1bhTJBsYE
mDMUizicDN%2BuutHXZBVLnsKYFnnOwO3k7wmieKIAtW5HkN672Hg%2Fz06eYL9Y%2
FvVQ5%2FZPpdJ1%2B67VQBZ4Po7KpHDli8PzKfMy%2BkZV1M0Oak7jblMutKwSpoYPb9q
YbXW1q4BchclrBenZcv%2B0Yc8yfJI38IA5ykRwmUSIPuanvooc14rAXITgKqld10jtkekPTJ%
2BmBJKs5qpfnPPVpH%2BNqnvgoNrQ%2BznnmkDHkMIDKE4MBQgX9FNEXhVj3xBReYR
YtS7eE0ddBbPqlgUzM%3D&X-Amz-Signature=6d588f3a6910ae38d4b4511d1143d54a077f
868ef40c3a91f4fa450483152bea&X-Amz-SignedHeaders=host&response-content-dispositio
n=inline